

Databases!

Introduction

DDL

DML

Data-Definition Language

Data-Manipulation Language

CREATE, DROP, ALTER

INSERT, UPDATE, DELETE, SELECT

Relation schema

instance

Database

Name of table + attributes + types

Records in the table at a certain point in time

a collection of Relations

Key

duplication is forbidden!

Primary key can fex. be 1D. Possibly Many keys, but 1PK

Foreign key refers to PK in other schema

Super key when a collection of columns are a key

Primitive types

all can be NULL

INT / INTEGER ← Round numbers

REAL / FLOAT ← Decimal numbers

CHAR(n) / VARCHAR(n) ← strings → fixed/variable length

DATE ← dates

Basic Queries

SELECT * FROM <relation> WHERE <condition>

LIKE matches string, '%' - any string '_' - any char
Distinct no duplicates

GROUP BY <columns> HAVING <condition>
ORDER BY <columns> ↑ only consider grouped records where...
↑ only for pretty outputs

<Relation 1> JOIN <Relation 2> ON <col 1> = <col 2>

SUM, AVG, COUNT, MIN, MAX,...
↑ aggregate functions

Complex Queries

Division

which x has ALL possible y

which x has ALL possible y
solution: group by x having count(y) = (subquery #Distinct y's)

Joins

INNER JOIN

A JOIN B ON ..

NATURAL-JOIN *i am dangerous joins*

CROSS JOIN rarely used

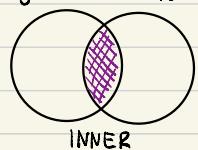
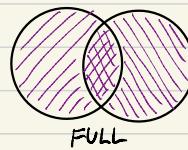
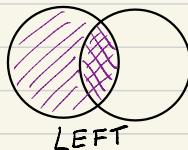
CROSS JOIN SELF JOIN

Can self-inhibit with 'down-regulating genes'

OUTER JOIN ← can be **LEFT, RIGHT, FULL**

OUTER JOIN ← can be LEFT, RIGHT, FULL

Left is main right is main Nothing removed



Set Operation

<Query 1> UNION / INTERSECT / EXCEPT <Query 2>

UNION we take 2 results, and plaster together

INTERSECT we take 2 results, and return common rows

INTERSECT we take 2 results, and return common rows
EXCEPT we return the ones from Q1 that are not in Q2

Duplicates are removed by default...

↳ you want 'em? solution: ALL

UNION ALL when Q1 & Q2 are disjoint!

EXISTS(Subquery) → TRUE if subquery non-empty
in these subqueries, we sometimes use attributes from main query

Views & Subqueries

Division

find x that have $\underbrace{\text{all } y}_{(\text{count}(*)) \text{ from } y}$ \leftarrow subquery!

Group by x Having $\text{count}(\text{distinct } y) = ()$

Subqueries

SELECT - FROM - WHERE as conditions / tables

Where $x = (\text{subquery})$ \leftarrow single tuple

Where $x \in (\text{subquery})$ \leftarrow multi-value

Where $\text{EXISTS } (\text{subquery})$ \leftarrow returns anything?
only this can include outer tables

Where $x \geq \text{ANY } (\text{subquery})$ \leftarrow multi-value, True for at least 1

$x \geq \text{ALL } (\text{subquery})$ \leftarrow multi-value, True for all

Views

CREATE VIEW <NAME>(<columns>) AS
<Query>;

SQL Programming

Functions

Callable, runs queries / commands + logic if & loops

How to use

From Java via JDBC OR SQL-script
May be faster than executing on a client...

Syntax:

```
DROP FUNCTION IF EXISTS function_name();  
CREATE FUNCTION function_name(IN var dtype) RETURNS dtype AS $$  
DECLARE var dtype;  
BEGIN Begin SQL put in local variable  
    SELECT something INTO var } Query  
    FROM Something...;  
    RETURN var; Return value  
END $$ LANGUAGE plpgsql;  
End Function text Specify Language
```

Triggers

Helps maintain integrity of data

Trigger: Insert / Update / Delete

Trigger-data: NEW. attribute AND/OR OLD. --

TG-OP variable containing operation type

1. Drop function & Trigger if exists
2. Create function 'returns TRIGGER', 'RETURN NEW'
3. Create Trigger ;

CREATE TRIGGER trigger-name

AFTER/BEFORE INSERT OR UPDATE OR DELETE
ON Relation

FOR EACH ROW EXECUTE PROCEDURE function

Transactions

When we want to make sure to do all the necessary code in 1 block!
all or nothing execution!

BEGIN; → if we make it to commit, then
COMMIT; ← it is done
ROLLBACK; ↗ Rollback then can undo what
, we just did. nothing is changed!

SAVEPOINT <name>;

ROLLBACK TO SAVEPOINT <name>;

SQL-Injections on web services
can delete your ENTIRE database

Entity Relationships

ER-diagrams

Why?

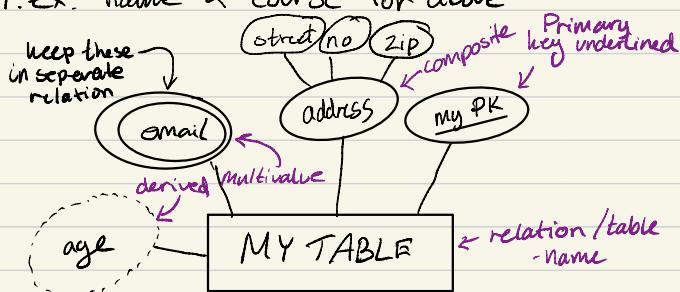
Way more interpretable!
Gives an overview of the DB schema

Entity Type

Corresponds to a record / instance relation name
F.ex. 'Students' with 'John', 'CSE305' in it

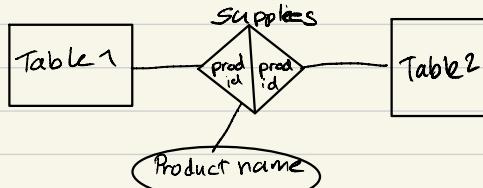
Attribute Type

Corresponds to an aspect attribute names
F.ex. name & course for above



Relations

Relate entities with roles



Cardinality

Min: 0 or 1

Max: 1, N/M/L/*...

Review Cardinality, 0...N

What is total and disjoint

& Generalization/Complexity

Star-schema \leftrightarrow

Practical Normalization

why?

Improving the design & runtime

Redundancy Issues

Functional Dependency (FD)

Solution

How?

e.g. zip \leftrightarrow city (have one, can deduce other)

i.e. we don't need both columns

Decomposition:

new table that has zip \rightarrow city mapping

1. create new table
2. insert mapping from prior table
3. remove column from prior

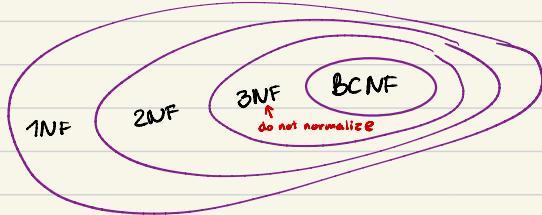
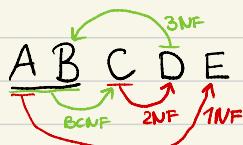
Normal Forms

NF Before we decompose; i.e. redundant cols.

BCNF After we decompose; i.e. redundancy removed

Boyce-Codd Normal Form

every attribute is dependent ONLY on the whole key



1NF

2NF

3NF

Part of key refers to other entity $\underline{AB} \underline{CD} \rightarrow \underline{ABC} \underline{AD}$

non-key entity refers to other entity $\underline{AB} \underline{CD} \rightarrow \underline{ABC} \underline{CD}$

non-key entity refers to key entity $\underline{AB} \underline{CD}$

Decompose 2NF before 1NF

Storage Hierarchy, Multicore OS

Cache
Main Memory
Local Disk
Remote Disk
Archival



Sequential access is faster than random access
when fetching:

from Disk : 4KB Blocks

from Main Memory : 64B Blocks

from caches : 4-8B Blocks

! Improve Locality for frequently accessed data

Parallelism / Multi-core

CPU Central Processing Unit

Multi-core several cores grouped

Multi-socket several layers grouped

Multi-core parallelism has to be actively exploited

Operating Systems

Resource Management → Hardware is limited

Indexing

Runtime Analysis :
CREATE INDEX

EXPLAIN ANALYZE <QUERY>

CREATE INDEX <name> ON <table> (<column>)

Goal: Avoid reading in random order.

Return 80% <

Point Queries
Range Queries

Full Table Scan: we read through EVERY record and return relevant tuples

<column> = <value>: easy if sorted by <column>

<column> BETWEEN <value1> and <value2>: easy if sorted by <column>

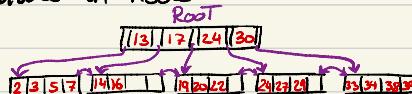
Solution: Indexes a sorted hashmap of a column's values

Built as hashing or Tree Search
point queries only

B+ trees

most common indexing type → Binary search tree
Dynamic, Balanced

4 values in node



Insertion to full node: split and move middle key to parent

Clustered

Data is stored according to the index...
need to fetch fewer blocks from memory
Only 1 clustered index per table

Un- Clustered

Should only be used for VERY selective queries

Covering

Contains all the information for the query...
i.e. count(*) or select <column> where <column> is part of index

Multi-index Firstly sorted by <attr1>, and secondly by <attr2>
good for point on <attr1> combined with range on <attr2>

Query Processing

Indexes

if covering: use

if clustered: use

if non-clustered: use ONLY when high selectivity

Multi-condition

use 2 non-clustered indices return intersect/union

and or
if 1 condition only returns few rows, scan 1 index, filter the "few"

use composite (& covering) indexing, return \cup

if remaining columns are few (Because takes up memory)

Join
evaluation
only 1st in
MySQL ↪

Approach 1: Nested loop

for tuple in R
for tuple in S
if r.ID = s.ID
output(r,s)

$O(R| + S|)$

2 x full table scans

Approach 2: Nested loop on index

binary $O((R| \cdot log(S|))$

easier to search index

Approach 3: Merge Join; sort both on join column, then
merge the results!

works well for clustered & covering index

$O((R| \cdot log(R|) + S| \cdot log(S|) + R| + S|)$

Sorting

merging

Approach 4: Hash Join; create hash table for 1 table, then
scan through using hash function on other table to join

works well when
fits in RAM

$O(R| + S|)$

Apply non-join conditions before hashing

join cost + RAM

for min(R|, S|)

Main memory :

MMDBS, NoSQL & Scaling Out

OLTP vs. OLAP

- only TRANSACTION processing vs. only ANALYTICS processing
- * short, simple Ts
 - * Many users!
 - * complex queries
 - * few HEAVY users

Optimized storage

Row-wise!

Column-wise!

Here we can use MMDBM
if all data fits in RAM ...
less I/O waiting time

MMDBMS
for OLTP

All data is kept in Main Memory... No waiting for disk I/O
Don't have to optimize for disk, BUT better cache utilization
We need to reduce log size on disk...

Scaling Up

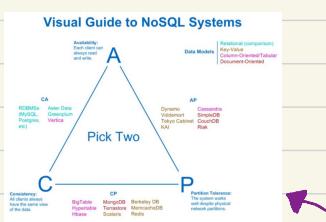
to keep performance running decently, we distribute users on more web servers
↳ a single server has an upper bound on performance

NoSQL

Usually

- ↳ Not Relational, Schema-less structure
- ↳ Not only SQL (often SQL-Like)
- ↳ Distributed Architecture*, ACID compliant **
- ↳ Mostly Open Source

Key-Value Stores like a dict ↳ flex URL → HTML
Document Stores like JSON ↳
Graph Stores like nodes, edges & properties ↳



* Distributed → consistent hashing - virtual server concept
** CAP Theorem → Workload sharing, failure handling, balance
True, but is about trade-off when you have a partitioned network!

Big Data Management

Big Data

Data Science & Real Life Application
How are big amounts of data used/saved?
Particle Accelerator, DNA, Video Surveillance and so on...

Big Data is a collection of processes

- ↳ DBs, Docs, feeds, Analog
 - collect, clean, ETL

The Three V's (Only Big Data if all are present)

- ↳ Volume! (amount)
- ↳ Velocity! (processing speed, accumulation speed)
data generation + processing action!
- ↳ Variety! (different ^{semi-structured!} types and sources!)
+ Veracity! (know whole data! not just parts of it...)
- ↳ Value! (data provides monetary / societal value!)

Clustering Big Data

Identify groups in data (f.ex. k-means)

Classifying Big Data

Find Rules to predict outcomes (f.ex. Decision Trees)

Access Patterns

Distributed Storage + Large Sequential Scans
Run complex processing Pipelines!

RDBMS is not relevant to big data...

↳ no support for unstructured

↳ little support for distributed processing

Same for NoSQL options

MapReduce

1. map Split input, 2. shuffle, 3. reduce, 4. gather

Example: Apache Hadoop (But! ^{complex mapping} lots of disk writes)

Spark

+ Spark SQL

Transform one 'Resilient Distributed Dataset' (RDD)
to another via operators...

Deep Pipelines + worker's memory sharing

